

CPSC 8810 - Deep Learning
Deep Learning Model to Detect Cyberbully
Actions in Images

Submitted By:

Vivek Koodli Udupa - (C12768888)

Shashi Shivaraju - (C88650674)

Clemson University

May 1, 2019

Abstract

This report explains the process involved in implementing a deep CNN and Faster-RCNN model in order to classify various cyberbully actions in an image and to detect the predator and victim in the same image.

1 Introduction

This report considers the problem of detection and classification of cyberbully actions and to identify the predator and victim in a bullying image using Deep learning models.

Cyberbullying is bullying that takes place over digital devices like cell phones, computers, and tablets. Cyberbullying can occur through SMS, Text, and apps, or online in social media, forums, or gaming where people can view, participate in, or share content. Cyberbullying includes sending, posting, or sharing negative, harmful, false, or mean content about someone else. Some cyberbullying crosses the line into unlawful or criminal behavior. With the prevalence of social media and digital forums, comments, photos, posts, and content shared by individuals can often be viewed by strangers as well as acquaintances. The content an individual shares online – both their personal content as well as any negative, mean, or hurtful content – creates a kind of permanent public record of their views, activities, and behavior. This public record can be thought of as an online reputation, which may be accessible to schools, employers, colleges, clubs, and others who may be researching an individual now or in the future. Cyberbullying can harm the online reputations of everyone involved – not just the person being bullied, but those doing the bullying or participating in it.[1]

Deep learning is a subset of machine learning where artificial neural networks, algorithms inspired by the human brain, learn from large amounts of data. The term ‘deep learning’ because the neural networks have various (deep) layers that enable learning. Deep learning allows machines to solve complex problems even when using a data set that is very diverse, unstructured and inter-connected. The more deep learning algorithms learn, the better they perform. [2]

This report describes modeling of a convolutional neural network for detecting and classifying cyberbully actions for a given image along with a RCNN model inorder to identify predator and victim present in that image. The cyberbullying actions considered in this project are laughing, pulling-hair, quarrel, slapping, punching, stabbing, gossiping, strangle and isolation. The CNN is trained using the provided image dataset which contain above mentioned 9 categories of cyberbully actions in them. The proposed RCNN is trained using provided images with ground truth bounding boxes for predator and victim classes.

2 Methods

In deep learning, a convolutional neural network (CNN, or ConvNet) is a class of deep neural networks, most commonly applied to analyzing visual imagery.[3] A convolutional neural network consists of an input and an output layer, as well as multiple hidden layers. The hidden layers of a CNN typically consist of convolutional layers, ReLU layer i.e. activation function, pooling layers, fully connected layers and normalization layers[4], which have been described in detail in section 2.1.

The implemented CNN model classifies the given input image into one of the nine categories of cyberbully actions. In order to further identify the predator and victim in the classified cyberbully image, an object detection model must be used. This can be achieved by using a Region based Convolutional Neural Network(RCNN).

In R-CNN, the CNN is forced to focus on a single region at a time because it minimizes the interference and it is expected that only a single object of interest will dominate in a given region. These regions in the R-CNN are detected by selective search algorithm followed by resizing so that the regions are of equal size before they are fed to a CNN for classification and bounding box regression. The drawback of RCNN is that selective search algorithm is computation expensive. Thus in order to overcome this drawback a Faster RCNN(F-RCNN) model is proposed.

F-RCNN approach is similar to the R-CNN algorithm but instead of feeding the region proposals to the CNN, the whole input image is fed to the CNN to generate a convolutional feature map. From the convolutional feature map, region of proposals are identified and warped into squares(Bounding Boxes) and by using a RoI(Region of Interest) pooling layer, region proposals are reshaped into fixed sizes so that it can be fed into a fully connected layer. Class of the proposed region and the offset of the bounding boxes are calculated by performing softmax on the RoI feature vector(Final layer of fully connected network)[6]

The implementation of our models based on CNN and FRCNN algorithm using an open source machine learning library PyTorch is described in the below sections.

2.1 Implementation of CNN Model

Our CNN model for classification consists of the following layers:

1. **Convolutional Layer:** The convolutional layer is the core building block of a CNN. The layer's parameters consist of a set of learnable filters (or kernels). During the forward pass, each filter is convolved across the width and height of the input volume, computing the dot product between the entries of the filter and the input and producing a 2-dimensional activation map of that filter. As a result, the network learns filters that activate when it detects some specific type of feature at some spatial position in the input.

2. **ReLU layer:** ReLU is the abbreviation of rectified linear unit, which applies the non-saturating activation function

$$f(x) = \max(0, x) \quad (1)$$

It effectively removes negative values from an activation map by setting them to zero. It increases the nonlinear properties of the decision function and of the overall network without affecting the receptive fields of the convolution layer.

3. **Max Pooling:** Another important concept of CNNs is pooling, which is a form of non-linear down-sampling. Max pooling is the most common non-linear function for down-sampling. It partitions the input image into a set of non-overlapping rectangles and, for each such sub-region, outputs the maximum.
4. **Fully Connected Layer:** Finally, after several convolutional and max pooling layers, the high-level reasoning in the neural network is done via fully connected layers. Neurons in a fully connected layer have connections to all activations in the previous layer, as seen in regular (non-convolutional) artificial neural networks. Their activations can thus be computed as an affine transformation, with matrix multiplication followed by a bias offset.
5. **Dropout Layer:** A single model can be used to simulate having a large number of different network architectures by randomly dropping out nodes during training. This is called dropout and offers a very computationally cheap and remarkably effective regularization method to reduce over fitting and generalization error in deep neural networks.[5]

The implementation details of our model is as follows:

1. **Image Pre-Processing:** The given image is resized to (256 x 256) pixels. Then the images are randomly flipped or rotated for the purpose of data augmentation. Then it is converted to a PyTorch tensor image and its values are normalized with a mean of 0.5 and Standard Deviation of 0.5.
2. **Convolution Layer 1_1:** The input to this layer is a preprocessed 3 channel tensor image from the previous layer. This layer performs 2D convolution using a 3x3 kernel with stride set to 1 and padding enabled to produce an output which is a 64 channel feature map. Xavier initialization is used to initialize the weights of this layer.
3. **ReLU Layer 1:** This layer applies a relu activation function to the 16 channel feature map.
4. **Convolution Layer 1_2 and ReLU:** The input to this layer is a 64 channel feature map from the previous layer. This layer performs 2D convolution using a 3x3 kernel with stride set to 1 and padding enabled to produce an output which is a 64 channel feature map. The output is normalized using ReLU.
5. **Batch Normalization:** This layer performs batch normalization inorder to avoid overfitting.
6. **Max Pooling Layer 1:** This layer down-samples the 256 x 256 64 channel feature map to 128 x 128 64 channel feature map.
7. **Convolution Layer 2_1 and ReLU:** The input to this layer is the 64 channel 128 x 128 feature map from the previous layer. This layer performs 2D convolution using a 3 x 3 kernel with stride set to 1 and padding enabled to produce an output which is a 128 channel feature map. ReLU activation is used to normalize the outputs.
8. **Convolution Layer 2_2 and ReLU:** The input to this layer is the 128 channel 128 x 128 feature map from the previous layer. This layer performs 2D convolution using a 3 x 3 kernel with stride set to 1 and padding enabled to produce an output which is a 32 channel feature map. ReLU activation is used to normalize the outputs.
9. **Batch Normalization:** This layer performs batch normalization inorder to avoid overfitting.
10. **Max Pooling Layer 2:** This layer down-samples the 128 x 128 128 channel feature map to 64 x 64 128 channel feature map.

11. **Convolution 3 with ReLU:** Similar to above convolution layers, Convolution 3 has 3 convolution layers, Convolution Layer 3_1, Convolution Layer 3_2 and Convolution Layer 3_3. The output of this layer is 64 x 64 512 channel feature map.
12. **Max Pooling Layer 3:** This layer down-samples the 64 x 64 512 channel feature map to 32 x 32 512 channel feature map.
13. **Convolution 4 and Convolution 5 with ReLU:** These two layers have similar configurations and structure as the Convolution 3 layer. The final output is a 8 x 8 512 channel feature map.
14. **Flattening Layer:** This layer flattens the 2D feature map to 1D feature map.
15. **Dropout Layer:** This layer randomly zeros some of the element of the input tensor with probability 0.4.
16. **Fully Connected Layer 1 and Relu layer 3:** This layer maps the 1D feature map into 4096 neurons.
17. **Fully Connected Layer 2 and Relu layer 4:** This layer maps the 4096 neurons to 1000 neurons.
18. **Fully Connected Layer 3 and softmax:** This layer maps 1000 neurons into 10 categories of classification and softmax for normalization.

2.2 Training the CNN Model

To train a deep learning model, the following parameters are considered:

1. **Epoch:** An epoch describes the number of times the algorithm sees the entire data set. So, each time the algorithm has seen all samples in the dataset, an epoch has completed.
2. **Batch Size:** The total number of training examples present in a single batch, wherein a batch is a subset of the entire data set.
3. **Iteration:** The number of batches needed to complete one Epoch.
4. **Learning Rate:** The learning rate or step size in machine learning is hyper-parameter which determines to what extent newly acquired information overrides old information.

The implemented model is trained using the below mentioned configuration:

1. Epoch = 100
2. Batch Size = 10
3. Learning Rate = 0.001

The model is trained with the given training dataset as per the below mentioned algorithm:

1. Initialize the CNN model with default parameters.
2. Create an instance of Adam optimizer for setting the learning rate
3. Create an instance of cross entropy loss
4. Initialize the optimizer with zero gradients
5. Feed a training input image from the current batch to the model to perform forward propagation
6. After the completion of forward propagation, calculate the cross entropy loss
7. Perform back propagation to minimize the loss
8. Update gradients
9. Iterate through step 4 for all the batches in the training dataset
10. Repeat the above steps for the given number of epochs
11. Save the trained model for testing purpose

2.3 Implementation of FRCNN Model

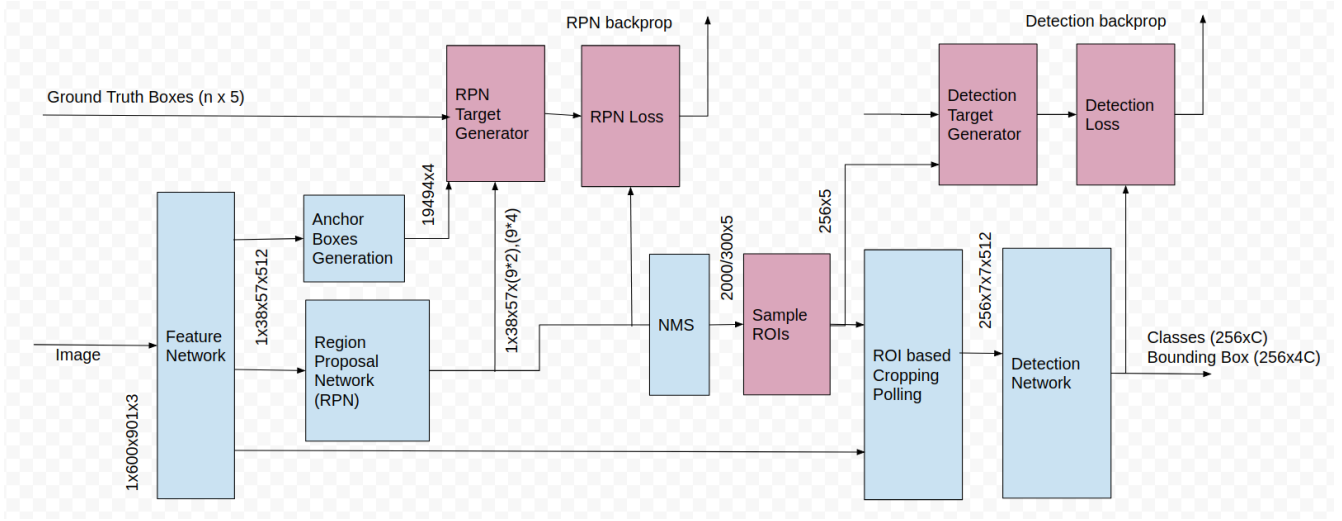


Figure 1: Faster-RCNN block diagram. The magenta colored blocks are active only during training. The numbers indicate size of the tensors[7].

The architecture for the object detection model, FRCNN is shown in Figure 1. To detect objects in the given image, the Faster RCNN uses two models which are RPN for generating region proposals and another detection which uses generated proposals to identify objects[10]. The basic building blocks of FRCNN are explained in detail below:

Region Proposal Network(RPN):

1. In the first step, the input image goes through a CNN which will output a set of convolutional feature maps.




Figure 2: CNN layer with the feature map output

2. In step 2, a sliding window is run spatially on these feature maps. The size of sliding window is $n \times n$ (generally 3×3). For each sliding window, a set of 9 anchors are generated which all

have the same center (x_a, y_a) but with 3 different aspect ratios and 3 different scales as shown below. Note that all these coordinates are computed with respect to the original image.

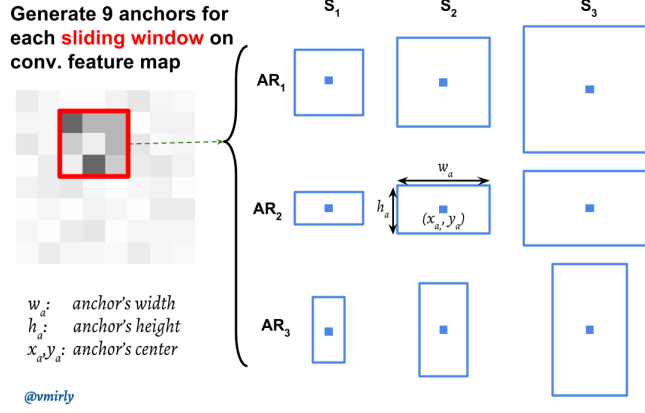


Figure 3: Anchors

For each anchor, a value p^* is computed which indicates how much the anchor overlaps with the groundtruth boxes. p^* is 1 if $\text{IoU} > 0.7$, -1 if $\text{IoU} < 0.3$ and 0 otherwise. where IoU is the intersection over union which is defined as: $\text{IoU} = \frac{\text{Anchor} \cap \text{GT Box}}{\text{Anchor} \cup \text{GT Box}}$

3. In step 3, Finally, the 3×3 spatial features extracted from those convolution feature maps (shown above within red box) are fed to a smaller network which has two tasks: classification (cls) and regression (reg). The output of regressor determines a predicted bounding-box (x, y, w, h), The output of classification sub-network is a probability p indicating whether the predicted box contains an object (1) or it is from background (0 for no object).

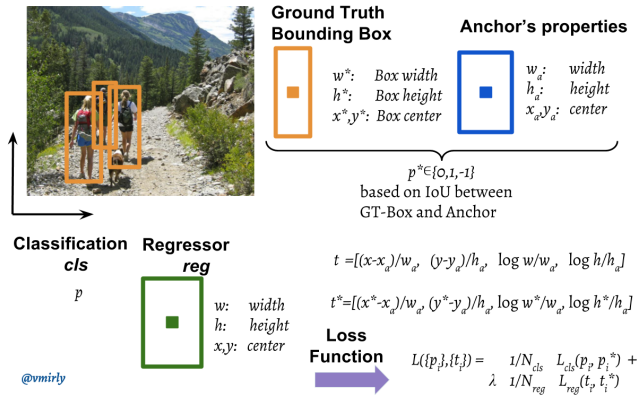


Figure 4: Classification and Regression

Non Maximum Supression(NMS): It is the process in which we remove/merge extremely

highly overlapping bounding boxes. The general idea of non-maximum suppression is to reduce the number of detections in a frame to the actual number of objects present. If the object in the frame is fairly large and more than 2000 object proposals have been generated, it is quite likely that some of these will have significant overlap with each other and the object. The pseudo code to implement NMS is given below:

```
- Take all the roi boxes [roi_array]
- Find the areas of all the boxes [roi_area]
- Take the indexes of order the probability score in descending order [order_array]
keep = []
while order_array.size > 0:
    - take the first element in order_array and append that to keep
    - Find the area with all other boxes
    - Find the index of all the boxes which have high overlap with this box
    - Remove them from order array
    - Iterate this till we get the order_size to zero (while loop)
- Output the keep variable which tells what indexes to consider.
```




Figure 5: An example of NMS

RoI Pooling: Region of interest pooling (also known as RoI pooling) purpose is to perform max pooling on inputs of non-uniform sizes to obtain fixed-size feature maps. RoI Pooling is done in three steps:

1. Dividing the region proposal into equal-sized sections (the number of which is the same as the dimension of the output)

2. Finding the largest value in each section
3. Copying these max values to the output buffer

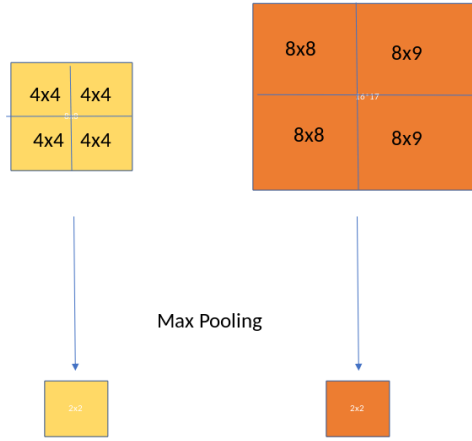


Figure 6: An example of ROI pooling

Detection Network: Using the region proposals generated by the RPN network, Fast R-CNN detection network is used to classify and regresses the bounding boxes. Here, ROI pooling is performed first and then the pooled area goes through CNN and two FC branches for class softmax and bounding box regressor[11].




Figure 7: Fast R-CNN detection Network

2.4 Training the FRCNN

The steps taken to develop FRCNN is as follows[8]:

1. Pre-train a CNN network on image classification tasks as described in section 2.2
2. Fine-tune the RPN (Region Proposal Network) end-to-end for the region proposal task, which is initialized by the pre-train image classifier. Positive samples have IoU (Intersection-over-Union) > 0.7 , while negative samples have IoU < 0.3 .

- (a) Slide a small $n \times n$ spatial window over the convolution feature map of the entire image.
- (b) At the center of each sliding window, we predict multiple regions of various scales and ratios simultaneously. An anchor is a combination of (sliding window center, scale, ratio).
3. Train a Fast R-CNN object detection model using the proposals generated by the current RPN
4. Then use the Fast R-CNN network to initialize RPN training. While keeping the shared convolutional layers, only fine-tune the RPN-specific layers. At this stage, RPN and the detection network have shared convolutional layers.
5. Finally fine-tune the unique layers of Fast R-CNN
6. Step 4-5 can be repeated to train RPN and Fast R-CNN alternatively if needed.

Loss functions for the FRCNN is calculated as follows:

$$L_{box}(\{p_i\}, \{t_i\}) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, p_i^*) + \frac{\lambda}{N_{box}} \sum_i p_i^* \cdot L_i^{smoott}(t_i - t_i^*) \quad (2)$$

$$L_{cls}(p_i, p_i^*) = -p_i^* \log p_i - (1 - p_i^*) \log(1 - p_i) \quad (3)$$

$$L = L_{cls} + L_{box} \quad (4)$$

where

p_i = Predicted probability of anchor i being an object

p_i^* = Ground truth label (binary) of whether anchor i is an object

t_i = Predicted four parameterized coordinates

t_i^* = Ground truth coordinates

N_{cls} = Normalization term, set to be mini-batch size (256) in the paper

N_{box} = Normalization term, set to the number of anchor locations (2400) in the paper

λ = A balancing parameter, set to be 10 in the paper

Please refer the appendix for the python implementation of the above described model.

3 Expected Results

Upon implementation and training of the model described in Section 2.3 the following results are expected.



Figure 8: Test image from slapping category



Figure 9: Result with bounding boxes for predator and victim

The test image shown in Figure 8 will be classified as cyberbully action : slapping with predator identified with red bounding box and the victim identified with green bounding box as shown in Figure 9.



Figure 10: Test image from slapping category



Figure 11: Result with bounding boxes for predator and victim

The test image shown in Figure 10 will be classified as cyberbully action : punching with predator identified with red bounding box and the victim identified with green bounding box as shown in Figure 11.

4 Conclusion

The goal of this project was to develop a deep neural network that takes an image as input and categorizes it into one of the 10 below mentioned categories of bullying. Further it was desired that the identification network would identify the predator and victim in the images classified as cyberbullying.

1) Gossiping 2) Isolation 3) Laughing 4) Pulling hair 5) Punching 6) Quarrel 7) Strangle 8) Slapping 9) Stabbing 10) Non bullying

The neural network was developed using PyTorch. The training dataset for the classification network was made up of 2494 images belonging to the 10 above mentioned categories.

In order to avoid overfitting in the CNN model for classification network, batch normalization and image augmentation strategies were implemented.

Unfortunately, due to time constraints, we were unable to complete the implementation of the researched FRCNN model. Please refer to the appendix for the partial implementation code of FRCNN model. We believe that the above researched model would provide satisfactory result for the given project problem statement as shown in [9].

5 References

- [1] <https://www.stopbullying.gov/cyberbullying/what-is-it/index.html>
- [2] <https://www.forbes.com/sites/bernardmarr/2018/10/01/what-is-deep-learning-ai-a-simple-guide-with-8-practical-examples/#434cffaa8d4b>
- [3] https://en.wikipedia.org/wiki/Convolutional_neural_network#Convolutional
- [4] <https://cs231n.github.io/convolutional-networks/>
- [5] <https://machinelearningmastery.com/dropout-for-regularizing-deep-neural-networks/>
- [6] <https://towardsdatascience.com/r-cnn-fast-r-cnn-faster-r-cnn-yolo-object-detection-al>
- [7] <https://medium.com/@whatdhack/a-deeper-look-at-how-faster-rcnn-works-84081284e1cd>
- [8] <https://lilianweng.github.io/lil-log/2017/12/31/object-recognition-for-dummies-part-3.html#fast-r-cnn>
- [9] <https://arxiv.org/abs/1506.01497>
- [10] <https://www.quora.com/How-does-the-region-proposal-network-RPN-in-Faster-R-CNN-work>
- [11] <https://towardsdatascience.com/review-faster-r-cnn-object-detection-f5685cb30202>

6 Appendix

6.1 Bully Detection CNN Model

```
1 # Importing libraries
2 import torch
3 import torchvision
4 import torchvision.transforms as transforms
5 import torchvision.datasets as datasets
6 import torch.nn as nn
7 import torch.nn.functional as F
8 from torch.autograd import Variable
9 import sys # For command Line arguments
10
11 #Defining the CNN
12 class CNNModel(nn.Module):
13     """ A CNN Model for image classification """
14
15     def __init__(self, image_size, op_size):
16         """ CNN layer to process the image """
17         super(CNNModel, self).__init__() # Super is used to refer to the base c
18
19         # Convolution Layer 1
20         self.cnn1_1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
21
22         # Xavier Initialization
23         nn.init.xavier_uniform_(self.cnn1_1.weight)
24
25         self.cnn1_2 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3,
26
27         # Batch Normalization
28         self.cnnBN1 = nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
29
30         # Max Pooling 1
31         self.maxpool1 = nn.MaxPool2d(kernel_size=2)
32
33         self.dropout = nn.Dropout(p=0.8)
34
35         # Convolution Layer 2
36         self.cnn2_1 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3
37         self.cnn2_2 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=
38         self.cnnBN2 = nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
39
40         # Max Pooling 2
41         self.maxpool2 = nn.MaxPool2d(kernel_size=2)
42
```

```

43     self.dropout = nn.Dropout(p=0.8)
44
45     # Convolution Layer 3
46     self.cnn3_1 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=
47     self.cnn3_2 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=
48     self.cnn3_3 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=
49     self.cnnBN3 = nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
50
51     # Max Pooling 3
52     self.maxpool3 = nn.MaxPool2d(kernel_size=2)
53
54     # Dropout Regularization
55     self.dropout = nn.Dropout(p=0.8)
56
57     # Convolution Layer 4
58     self.cnn4_1 = nn.Conv2d(in_channels=256, out_channels=512, kernel_size=
59     self.cnn4_2 = nn.Conv2d(in_channels=512, out_channels=512, kernel_size=
60     self.cnn4_3 = nn.Conv2d(in_channels=512, out_channels=512, kernel_size=
61     self.cnnBN4 = nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
62
63     # Max Pooling 4
64     self.maxpool4 = nn.MaxPool2d(kernel_size=2)
65
66     # Dropout Regularization
67     self.dropout = nn.Dropout(p=0.5)
68
69     # Convolution Layer 5
70     self.cnn5_1 = nn.Conv2d(in_channels=512, out_channels=512, kernel_size=
71     self.cnn5_2 = nn.Conv2d(in_channels=512, out_channels=512, kernel_size=
72     self.cnn5_3 = nn.Conv2d(in_channels=512, out_channels=512, kernel_size=
73     self.cnnBN5 = nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
74
75     # Max Pooling 5
76     self.maxpool5 = nn.MaxPool2d(kernel_size=2)
77
78     # Dropout Regularization
79     self.dropout = nn.Dropout(p=0.8)
80 #-----
81     # Fully connected linear layer
82     #self.fc1 = nn.Linear(32*75*75 , 9) #32 channels, 75x75 final image si
83     self.fc1 = nn.Linear(512*image_size*image_size, 4096)
84 #32 channels, 7x7 final image size
85     self.relu4 = nn.ReLU()
86
87     self.fc2 = nn.Linear(4096, 1000) #32 channels, 7x7 final image size
88     self.relu5 = nn.ReLU()

```

```

89         self.fc3 = nn.Linear(1000, 10)    #32 channels, 7x7 final image size
90
91
92
93     #Image size = 28x28 -> 13x13 after first pooling
94     #14x14 after padding = 1
95     #7x7 after second pooling
96
97     def forward(self, x):
98         """ Forward Propogation for classification """
99
100        #CNN layer 1
101        out = self.cnn1_1(x)
102        out = F.relu(out)
103        out = self.cnn1_2(out)
104        out = F.relu(out)
105        out = self.cnnBN1(out)
106        out = self.maxpool1(out)
107
108        #out = self.dropout(out)
109
110        #CNN layer 2
111        out = self.cnn2_1(out)
112        out = F.relu(out)
113        out = self.cnn2_2(out)
114        out = F.relu(out)
115        out = self.cnnBN2(out)
116        out = self.maxpool2(out)
117
118        #out = self.dropout(out)
119
120        #CNN layer 3
121        out = self.cnn3_1(out)
122        out = F.relu(out)
123        out = self.cnn3_2(out)
124        out = F.relu(out)
125        out = self.cnn3_3(out)
126        out = F.relu(out)
127        out = self.cnnBN3(out)
128        out = self.maxpool3(out)
129
130        #out = self.dropout(out)
131
132        #CNN layer 4
133        out = self.cnn4_1(out)
134        out = F.relu(out)
135        out = self.cnn4_2(out)

```

```

136         out = F.relu(out)
137         out = self.cnn4_3(out)
138         out = F.relu(out)
139         out = self.cnnBN4(out)
140         out = self.maxpool4(out)
141
142         #out = self.dropout(out)
143
144         #CNN layer 5
145         out = self.cnn5_1(out)
146         out = F.relu(out)
147         out = self.cnn5_2(out)
148         out = F.relu(out)
149         out = self.cnn5_3(out)
150         out = F.relu(out)
151         out = self.cnnBN5(out)
152         out = self.maxpool5(out)
153
154         out = self.dropout(out)
155
156         # Resize the tensor, -1 decides the best dimension automatically
157         #out = out.view(out.size(0), -1)
158         out = out.view(out.size(0), -1)
159
160         # Dropout
161         #out = self.dropout(out)
162
163         # Fully connected 1
164         out = self.fc1(out)
165         out = F.relu(out)
166
167         out = self.fc2(out)
168         out = F.relu(out)
169
170         out = self.fc3(out)
171
172         out = F.log_softmax(out, dim=0) #Softmax along Row
173         # Return
174         return out

```

6.2 Training Code

```

1 # Importing libraries
2 import torch
3 import torchvision
4 import torchvision.transforms as transforms
5 import torchvision.datasets as datasets

```

```

6 import torch.nn as nn
7 import torch.nn.functional as F
8 from torch.autograd import Variable
9 import sys # For command Line arguments
10 import os
11 from shutil import copyfile
12 from detection_Model import CNNModel
13
14 # Hyperparameter initialization
15 n_epoch          = 100
16 n_class          = 10
17 batch_size       = 10
18 learning_rate    = 0.0001
19
20 # check if GPU is available
21 print(torch.cuda.current_device())
22 print(torch.cuda.device(0))
23 print(torch.cuda.device_count())
24 print(torch.cuda.get_device_name(0))
25
26 #To run on GPU
27 device = torch.device("cuda:0")
28 dtype = torch.float
29 # Sorting out the data
30
31 # Image parameters
32 n_cnn = 5 #Number of CNN layer
33 img_size = (256,256)
34 conv_size = int( img_size[0]/ (2**n_cnn)) # image_size / 8 for 3 cnn layer.
35 train_img = "../TrainingData"
36 Model = "../Model"
37
38 # Define the transformation
39 transform = transforms.Compose( [transforms.Resize(img_size),
40                                transforms.RandomRotation((90, 360)),
41                                transforms.RandomVerticalFlip(),
42                                transforms.RandomHorizontalFlip(),
43                                transforms.ColorJitter(),
44                                transforms.ToTensor(),
45                                transforms.Normalize((0.5, 0.5, 0.5),(0.5, 0.5, 0.5))
46                                ])
47
48 # Training dataset
49 train_dataset = datasets.ImageFolder(root=train_img, transform=transform)
50
51 # Placing data into dataloader for better accessibility
52 # Shuffle training dataset to eliminate bias

```

```

53 train_loader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=ba
54
55 # Instance creation
56 #model = CNNModel(conv_size, n_class).cuda()
57 model = nn.DataParallel(CNNModel(conv_size, n_class))
58 model = model.to(device)
59 # Create instance of loss
60 criterion = nn.CrossEntropyLoss()
61
62 # Create instance of optimizer (Adam)
63 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
64
65 # Model Training
66
67 n_iteration = 0
68
69 for epoch in range(n_epoch):
70     total = 0
71     correct = 0
72     for i, (images, labels) in enumerate(train_loader):
73         # Wrap into Variable
74         #images = Variable(images).cuda()
75         #labels = Variable(labels).cuda()
76         images = Variable(images).to(device)
77         labels = Variable(labels).to(device)
78
79         # Clear the gradients
80         optimizer.zero_grad()
81
82         # Forward propogation
83         #outputs = model(images).cuda()
84         outputs = model(images)
85
86         # Loss calculation ( softmax )
87         loss = criterion(outputs, labels)
88
89         # Backpropogation
90         loss.backward()
91
92         # Update Gradients
93         optimizer.step()
94
95         n_iteration += 1
96
97         # Total number of labels
98         total += labels.size(0)
99

```

```

100         # obtain prediction from max value
101         _, predicted = torch.max(outputs.data, 1)
102
103         # Calculate the number of right answers
104         correct += (predicted == labels).sum().item()
105
106         # Print loss and accuracy
107         if (i + 1) % 10 == 0:
108             print('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}, Accuracy: {:.2f}%
109
110 # Saving the trained model
111 if not os.path.exists(Model):
112     os.makedirs(Model)
113 torch.save(model.state_dict(), "./Model/model.pth")
114 print("Model saved at ./Model/model.pth")

```

6.3 Test Code

```

1 # Importing libraries
2 import torch
3 import torchvision
4 import torchvision.transforms as transforms
5 import torchvision.datasets as datasets
6 import torch.nn as nn
7 import torch.nn.functional as F
8 from torch.autograd import Variable
9 import sys # For command Line arguments
10 import os
11 from shutil import copyfile
12 from detection_Model import CNNModel
13 from PIL import Image
14
15 # loading the input test image
16 if(len(sys.argv)<2):
17     sys.exit("Please specify an image to test")
18 else:
19     test_img_filename = sys.argv[1]
20
21 img_size = (256,256)
22 n_cnn = 3
23 conv_size = int(img_size[0] /(2**((n_cnn + 1)) )
24 test_img = "./TestData/test/"
25 test_img1 = "./TestData"
26 Model = "./Model"
27
28 device = torch.device("cuda:0")
29

```

```

30 # Hyperparameter initialization
31 batch_size      = 1
32
33 # Define the transformation
34 transform = transforms.Compose( [transforms.Resize(img_size),
35                                 transforms.ToTensor(),
36                                 transforms.Normalize((0.5, 0.5, 0.5),(0.5, 0.5, 0.5))]
37
38
39
40 # Testing dataset
41 test_dataset = Image.open(test_img_filename)
42 test_loader = transform(test_dataset)
43
44
45 # Image parameters
46 n_class      = 10
47
48 model = nn.DataParallel(CNNModel(conv_size, n_class))
49 model = model.to(device)
50 model.load_state_dict(torch.load('./Model/model.pth'))
51 model.eval().to(device)
52
53 def ten_to_str(x):
54     """ Function to convert tensor label to a string """
55     str_label = ["gossiping", "isolation", "laughing", "nonbullying", "pullingh
56     return str_label[x]
57
58 # Testing the model
59 with torch.no_grad():
60     images = Variable(test_loader, requires_grad=True)
61     images = images.unsqueeze(0)
62     images = images.to(device)
63     outputs = model(images)
64     _, predicted = torch.max(outputs.data, 1)
65     predicted = predicted.item()
66     print("{}".format(ten_to_str(predicted)))

```

6.4 FRCNN

```

1 # Imports
2 import torch
3 import torchvision
4 import torch.nn as nn
5 import numpy as np
6 import matplotlib.pyplot as plt
7 from torchvision import transforms

```



```

8  from PIL import Image
9
10 # Debug flag, set to 1 to get debug messages
11 __DEBUG__ = 0
12
13 image_size = (256, 256)
14 # Sample Black background image
15 image = torch.zeros((1, 3, image_size[0], image_size[1])).float()
16
17 if __DEBUG__:
18     print("size of image tensor: %s " %(image.size()))
19
20 # Generate Sample Bounding Box
21 bbox = torch.FloatTensor([[20, 30, 200, 150],[150, 200, 220, 250]]) #[Ymin, Xmin, Ymax, Xmax]
22 labels = torch.LongTensor([6, 8])
23 sub_sample = 16
24
25 # Squeeze to remove the first dimention of tensor (convert from 4d to 3d I think)
26 pil_image = transforms.ToPILImage()(image.squeeze())
27
28 #plt.imshow(pil_image)
29 if __DEBUG__:
30     pil_image.show()
31
32 # Create a dummy image
33 dummy_img = torch.zeros((1, 3, image_size[0], image_size[1])).float()
34
35 #print(dummy_img)
36
37 # Load vgg16
38 model = torchvision.models.vgg16(pretrained=True)
39
40 # List out all the features
41 fetrs = list(model.features)
42
43 # Pass dummy image through layers to check for layer whose output matches the req
44 req_fetrs = []
45 # Clone dummy image and pass it through all layers and check for layer output size
46 clone_img = dummy_img.clone()
47 for lyr in fetrs:
48     clone_img = lyr(clone_img)
49     if clone_img.size()[2] < 256//16:
50         break
51     req_fetrs.append(lyr)
52     out_channels = clone_img.size()[1]
53
54 if __DEBUG__:

```

```

55     print("Length of required features: ", len(req_fetrs))
56     print("Number of Out channels: ", out_channels)
57
58 # Convert required features into sequential module
59 frcnn_fe = nn.Sequential(*req_fetrs)
60
61 # Using frcnn as backend compute features for dummy image
62 out_map = frcnn_fe(image)
63
64 if __DEBUG__:
65     print("Out map size is: ", out_map)
66
67 # Creating Anchors
68 # -----
69 # Define 3 ratio and scales that we will be using
70 ratio = [0.5, 1, 2]
71 anchor_scales = [8, 16, 32]
72
73 # Number of Ratios and anchor scales
74 n_ratio = len(ratio)
75 n_scales = len(anchor_scales)
76
77 # Base for the anchor
78 anchor_base = np.zeros((n_ratio * n_scales, 4), dtype=np.float32)
79
80 if __DEBUG__:
81     print("anchor base is: \n", anchor_base)
82
83 # Define center for base anchor
84 center_y = sub_sample / 2.
85 center_x = sub_sample / 2.
86
87 if __DEBUG__:
88     print("Center for base anchor is: (%s, %s) " %(center_x, center_y))
89
90 # Generating Anchors for first feature map pixel
91 # Iterate through all ratios and scales
92 for i in range(n_ratio):
93     for j in range(n_scales):
94         h = sub_sample * anchor_scales[j] * np.sqrt(ratio[i])
95         w = sub_sample * anchor_scales[j] * np.sqrt(1. / ratio[i])
96
97         index = i * n_scales + j
98
99         anchor_base[index, 0] = center_y - h / 2. #y_min
100        anchor_base[index, 1] = center_x - w / 2. #x_min
101        anchor_base[index, 2] = center_y + h / 2. #y_max

```

```

102         anchor_base[index, 3] = center_x + w / 2. #x_max
103
104     if __DEBUG__:
105         print("anchor bases: \n", anchor_base)
106         print("Negative anchors represent the ones that are out of the image bounda
107
108     # Generationg anchors for all feature map pixels
109     feature_size = (image_size[0] // sub_sample)
110     # 16 sub_samples in feature map where each has dimension 16*16
111     center_x = np.arange(sub_sample, (feature_size + 1) * 16, 16)
112     center_y = np.arange(sub_sample, (feature_size + 1) * 16, 16)
113
114     # Generation Centers
115     center = np.zeros((len(center_x) * len(center_y), 2))
116     index = 0
117     for x in range(len(center_x)):
118         for y in range(len(center_y)):
119             center[index, 1] = center_x[x] - int(sub_sample / 2)
120             center[index, 0] = center_y[y] - int(sub_sample / 2)
121             index += 1
122
123     # Generating anchors for above generated centers
124     num_anchors_per_pixel = n_ratio * n_scales
125     anchors = np.zeros(((feature_size * feature_size * num_anchors_per_pixel), 4))
126
127     index = 0
128     for c in center:
129         center_y, center_x = c
130         for i in range(n_ratio):
131             for j in range(n_scales):
132                 h = sub_sample * anchor_scales[j] * np.sqrt(ratio[i])
133                 w = sub_sample * anchor_scales[j] * np.sqrt(1. / ratio[i])
134
135                 anchors[index, 0] = center_y - h / 2.
136                 anchors[index, 1] = center_x - w / 2.
137                 anchors[index, 2] = center_y + h / 2.
138                 anchors[index, 3] = center_x + w / 2.
139                 index += 1
140
141     if __DEBUG__:
142         print("Total anchors size is: ", anchors.shape)
143
144     # Labeling the anchors
145     # [Ymin, Xmin, Ymax, Xmax] format
146     bbox = np.asarray([[20, 30, 200, 150], [150, 200, 220, 250]], dtype=np.float32)
147     labels = np.asarray([6, 8])
148

```

```

149 # Find the index of the anchors that are inside the image boundary
150 index_inside = np.where(
151     (anchors[:, 0] >= 0) &
152     (anchors[:, 1] >= 0) &
153     (anchors[:, 2] <= image_size[1]) &
154     (anchors[:, 3] <= image_size[0])
155 ) [0]
156
157 if __DEBUG__:
158     print("anchors that are insdie the image are: \n", index_inside)
159     print("\nNumber of anchors inside the image boundary: ", index_inside.shape)
160
161 # Make a label array and fill it with -1
162 label = np.empty((len(index_inside), ), dtype=np.int32)
163 label.fill(-1)
164
165 if __DEBUG__:
166     print("Created Label size is %s and index inside size is %s" %(label.size,
167
168 # Array with valid anchor boxes
169 anchor_valid = anchors[index_inside]
170
171 if __DEBUG__:
172     print("Valid anchor box shape is: ", anchor_valid.shape)
173
174 # Calculate IoU for valid anchor boxes
175 ious = np.empty((len(anchor_valid), 2), dtype=np.float32)
176 if __DEBUG__:
177     print("Bounding Boxes are : \n", bbox)
178
179 for num1, i in enumerate(anchor_valid):
180     # ymin, xmin, ymax, xmax format for anchors
181     ya1, xa1, ya2, xa2 = i
182     # anchor area = height * width
183     area_anchor = (ya2 - ya1) * (xa2 - xa1)
184     for num2, j in enumerate(bbox):
185         yb1, xb1, yb2, xb2 = j
186         area_box = (yb2 - yb1) * (xb2 - xb1)
187
188         intersection_x1 = max([xb1, xa1])
189         intersection_y1 = max([yb1, ya1])
190         intersection_x2 = min([xb2, xa2])
191         intersection_y2 = min([yb2, ya2])
192
193         # Check for intersection
194         if (intersection_x1 < intersection_x2) and (intersection_y1 < intersect
195             area_intersection = (intersection_y2 - intersection_y1) * (intersec

```

```

196         #intersection over union
197         iou = area_intersection / (area_anchor + area_box - area_intersecti
198     else:
199         # In case of No overlap/ intersection
200         iou = 0.
201
202     ious[num1, num2] = iou
203
204 if __DEBUG__:
205     print("all the iou count: ", ious.shape)
206
207 # Case-1
208 # Highest IoU for each gt and corresponding anchor
209 # Location of max Iou
210 gt_argmax_ious = ious.argmax(axis=0)
211 if __DEBUG__:
212     print("Indices of MAX IoU: ", gt_argmax_ious)
213
214 # Value of Max IoU
215 gt_max_ious = ious[gt_argmax_ious, np.arange(ious.shape[1])]
216 if __DEBUG__:
217     print("Values of MAX IoU: ", gt_max_ious)
218
219 # Case-2
220 # Highest Iou In between every anchor
221 argmax_ious = ious.argmax(axis=1)
222 if __DEBUG__:
223     print("shape of argmax_ious: ", argmax_ious.shape)
224     print("MAX Iou indices for every anchor: ", argmax_ious)
225
226 max_ious = ious[np.arange(len(index_inside)), argmax_ious]
227 if __DEBUG__:
228     print("MAX IoU values: ", max_ious)
229
230 # Anchor Box that has the HIGHEST IoU with GT
231 gt_argmax_ious = np.where(ious == gt_max_ious)[0]
232 if __DEBUG__:
233     print("Ultimate MAX IoU: ", gt_argmax_ious)
234
235 # Assigning Labels which helps to compute Loss
236 # Defining thresholds
237 pos_thres = 0.7
238 neg_thres = 0.3
239
240 # Assign negative label (0) to all anchors that have IoU < 0.3
241 label[max_ious < neg_thres] = 0
242

```

```

243 # Assign positive label (1) to anchor boxes with highest IoU with GT
244 label[gt_argmax_iou] = 1
245
246 # Assign positive label (1) to anchor boxes with IoU > 0.7
247 label[max_iou >= pos_thres] = 1
248
249
250 # =====
251 #                                TRAINING RPN
252 # =====
253
254 # Define positive and negative anchor sample parameters
255 num_samples = 128
256 pos_ratio = 0.5
257 num_pos = pos_ratio * num_samples
258
259 # Picking Positive Samples
260 pos_index = np.where(label == 1)[0]
261
262 if len(pos_index) > num_pos:
263     disable_index = np.random.choice(pos_index, size=(len(pos_index) - num_pos))
264     label[disable_index] = -1
265
266 # Picking Negative Samples
267 num_neg = num_samples * np.sum(label == 1)
268 neg_index = np.where(label == 0)[0]
269
270 if len(neg_index) > num_neg:
271     disable_index = np.random.choice(neg_index, size=(len(neg_index) - num_neg))
272     label[disable_index] = -1
273
274 # GT with MAX IoU for each anchor
275 max_iou_bbox = bbox[argmax_iou]
276
277 # Convert [ymin, xmin, ymax, xmax] format to [center_y, center_x, h, w] format
278
279 height = anchor_valid[:, 2] - anchor_valid[:, 0]
280 width = anchor_valid[:, 3] - anchor_valid[:, 1]
281 center_y = anchor_valid[:, 0] + 0.5 * height
282 center_x = anchor_valid[:, 1] + 0.5 * width
283
284 base_height = max_iou_bbox[:, 2] - max_iou_bbox[:, 0]
285 base_width = max_iou_bbox[:, 3] - max_iou_bbox[:, 1]
286 base_center_y = max_iou_bbox[:, 0] + 0.5 * base_height
287 base_center_x = max_iou_bbox[:, 1] + 0.5 * base_width
288
289 # Find the locations

```

```

290 eps = np.finfo(height.dtype).eps
291 height = np.maximum(height, eps)
292 width = np.maximum(width, eps)
293
294 dy = (base_center_y - center_y) / height
295 dx = (base_center_x - center_x) / width
296 dh = np.log(base_height / height)
297 dw = np.log(base_width / width)
298
299 anchor_locs = np.vstack((dy, dx, dh, dw)).transpose()
300 if __DEBUG__:
301     print("Shape of anchor locations", anchor_locs.shape)
302     print("Anchor locations", anchor_locs)
303
304 # Final Labels
305 anchor_labels = np.empty((len(anchors),), dtype=label.dtype)
306 anchor_labels.fill(-1)
307 anchor_labels[index_inside] = label
308
309 # Final Locations
310 anchor_locations = np.empty((len(anchors),) + anchors.shape[1:], dtype=anchor_l
311 anchor_locations.fill(0)
312 anchor_locations[index_inside, :] = anchor_locs
313
314 # =====
315 #           Region Proposal Network
316 # =====
317
318 mid_channels = 512
319 in_channels = 512
320 n_anchor = 9
321
322 conv1 = nn.Conv2d(in_channels, mid_channels, 3, 1, 1)
323
324 # Bounding Box Regressor network
325 reg_layer = nn.Conv2d(mid_channels, n_anchor * 4, 1, 1, 0)
326
327 # Classifier network
328 cls_layer = nn.Conv2d(mid_channels, n_anchor * 2, 1, 1, 0)
329
330 # Initialization
331 # convolution sliding layer
332 conv1.weight.data.normal_(0, 0.01)
333 conv1.bias.data.zero_()
334
335 # Regression layer
336 reg_layer.weight.data.normal_(0, 0.01)

```

```

337 reg_layer.bias.data.zero_()
338
339 # classification layer
340 cls_layer.weight.data.normal_(0, 0.01)
341 cls_layer.bias.data.zero_()
342
343 # Training
344 # -----
345
346 x = conv1(out_map)
347 pred_anchor_locs = reg_layer(x)
348 pred_cls_scores = cls_layer(x)
349
350 if __DEBUG__:
351     print("predicted class score shape: ", pred_cls_scores.shape)
352     print("predicted anchor location shape: ", pred_anchor_locs.shape)
353
354 # Rearrange the tensors to align with anchor targets
355 pred_anchor_locs = pred_anchor_locs.permute(0, 2, 3, 1).contiguous().view(1, -1)
356 if __DEBUG__:
357     print("Rearranged anchor location shape: ", pred_anchor_locs.shape)
358
359 pred_cls_scores = pred_cls_scores.permute(0, 2, 3, 1).contiguous()
360 if __DEBUG__:
361     print("Rearranged class score shape: ", pred_cls_scores.shape)
362
363 # Calculate the Objectness score
364 objectness_score = pred_cls_scores.view(1, 16, 16, 9, 2)[: , : , : , : , 1].contiguous()
365 if __DEBUG__:
366     print("Objectness Score shape: ", objectness_score.shape)
367
368 pred_cls_scores = pred_cls_scores.view(1, -1, 2)
369 if __DEBUG__:
370     print("predicted class score shape final: ", pred_cls_scores.shape)
371
372 # =====
373 # Generationg Proposals
374 # =====
375
376 # Define parameters for training and testing
377
378 nms_thresh = 0.7 # Non- Maximum Supression Threshold
379 n_train_pre_nms = 12000 # number of bboxes before nms during training
380 n_train_post_nms = 2000 # number of bboxes after nms during training
381 n_test_pre_nms = 6000 # number of bboxes before nms during testing
382 n_test_post_nms = 300 # number of bboxes after nms during testing
383 min_size = 16 # minimum height of the object required to create a proposal

```



```

384
385 # Convert anchors from [ymin, xmin, ymax, xmax] to [center_Y, center_x, h, w] f
386 anc_height = anchors[:, 2] - anchors[:, 0]
387 anc_width = anchors[:, 3] - anchors[:, 1]
388 anc_ctr_y = anchors[:, 0] + 0.5 * anc_height
389 anc_ctr_x = anchors[:, 1] + 0.5 * anc_width
390
391 # Convert prediction locations and objectness score to numpy array
392 pred_anchor_locs_numpy = pred_anchor_locs[0].data.numpy()
393 objectness_score_numpy = objectness_score[0].data.numpy()
394
395 dy = pred_anchor_locs_numpy[:, 0::4]
396 dx = pred_anchor_locs_numpy[:, 1::4]
397 dh = pred_anchor_locs_numpy[:, 2::4]
398 dw = pred_anchor_locs_numpy[:, 3::4]
399
400 ctr_y = dy * anc_height[:, np.newaxis] + anc_ctr_y[:, np.newaxis]
401 ctr_x = dx * anc_width[:, np.newaxis] + anc_ctr_x[:, np.newaxis]
402 h = np.exp(dh) * anc_height[:, np.newaxis]
403 w = np.exp(dw) * anc_width[:, np.newaxis]
404
405 # Region of Interest
406 roi = np.zeros(pred_anchor_locs_numpy.shape, dtype=pred_anchor_locs_numpy.dtype)
407 roi[:, 0::4] = ctr_y - 0.5 * h
408 roi[:, 1::4] = ctr_x - 0.5 * w
409 roi[:, 2::4] = ctr_y + 0.5 * h
410 roi[:, 3::4] = ctr_x + 0.5 * w
411
412 # Clip the predicted boxes to the image size
413 roi[:, slice(0, 4, 2)] = np.clip(roi[:, slice(0, 4, 2)], 0, image_size[0])
414 roi[:, slice(1, 4, 2)] = np.clip(roi[:, slice(1, 4, 2)], 0, image_size[1])
415
416 if __DEBUG__:
417     print("Region of interest: ", roi)
418
419 # Remove predicted boxes with either height of width < threshold
420 hs = roi[:, 2] - roi[:, 0]
421 ws = roi[:, 3] - roi[:, 1]
422 keep = np.where((hs >= min_size) & (ws >= min_size))[0]
423 roi = roi[keep, :]
424 score = objectness_score_numpy[keep]
425
426 # Sort (proposal, score) in descending order
427 order = score.ravel().argsort()[::-1]
428
429 # Take top values?
430 order = order[:n_train_pre_nms]

```

```

431 roi = roi[order, :]
432
433 # Calculate Region Proposals
434 y1 = roi[:, 0]
435 x1 = roi[:, 1]
436 y2 = roi[:, 2]
437 x2 = roi[:, 3]
438
439 area = (x2 - x1 + 1) * (y2 - y1 + 1)
440 order = score.argsort()[::-1]
441
442 keep = []
443
444 while order.size > 0:
445     i = order[0]
446     keep.append(i)
447     xx1 = np.maximum(x1[i], x1[order[1:]])
448     yy1 = np.maximum(y1[i], y1[order[1:]])
449     xx2 = np.minimum(x2[i], x2[order[1:]])
450     yy2 = np.minimum(y2[i], y2[order[1:]])
451
452     w = np.maximum(0.0, xx2 - xx1 + 1)
453     h = np.maximum(0.0, yy2 - yy1 + 1)
454     inter = w * h
455
456     ovr = inter / (area[i] + area[order[1:]] - inter)
457
458     inds = np.where(ovr <= nms_thresh)[0]
459     order = order[inds + 1]
460
461 keep = keep[:n_train_post_nms]
462 roi = roi[keep]
463
464 if __DEBUG__:
465     print("Final region proposal count: ", roi.shape)
466
467 # Proposal Targets
468 n_sample = 128
469 pos_ratio = 0.25
470 pos_iou_thresh = 0.5
471 neg_iou_thresh_hi = 0.5
472 neg_iou_thresh_lo = 0.0
473
474 # Calculate IoU
475 ious = np.empty((len(roi), 2), dtype=np.float32)
476 ious.fill(0)
477 for num1, i in enumerate(roi):

```

```

478     ya1, xa1, ya2, xa2 = i
479     anchor_area = (ya2 - ya1) * (xa2 - xa1)
480     for num2, j in enumerate(bbox):
481         yb1, xb1, yb2, xb2 = j
482         box_area = (yb2 - yb1) * (xb2 - xb1)
483
484         inter_x1 = max([xb1, xa1])
485         inter_y1 = max([yb1, ya1])
486         inter_x2 = min([xb2, xa2])
487         inter_y2 = min([yb2, ya2])
488
489         if(inter_x1 < inter_x2) and (inter_y1 < inter_y2):
490             inter_area = (inter_y2 - inter_y1) * (inter_x2 - inter_x1)
491             iou = inter_area / (anchor_area + box_area - inter_area)
492         else:
493             iou = 0.
494
495         ious[num1, num2] = iou
496
497 if __DEBUG__:
498     print("Proposal Targets: ", ious.shape)
499
500 # GT with max IoU for each region
501 gt_assignment = ious.argmax(axis=1)
502 max_iou = ious.max(axis=1)
503
504 if __DEBUG__:
505     print("GT location with max IoU for each region and max IoU's")
506     print(gt_assignment)
507     print(max_iou)
508
509 # Assign label to each proposal
510 gt_roi_label = labels[gt_assignment]
511
512 if __DEBUG__:
513     print("GT labels: ", gt_roi_label)

```